

# 第8章：索引结构

## Index Structures

邹兆年

哈尔滨工业大学  
计算机科学与技术学院  
海量数据计算研究中心  
电子邮件: [znzou@hit.edu.cn](mailto:znzou@hit.edu.cn)

2021年春

## Outline<sup>1</sup>

- ① Hash-based Index Structures
  - Extensible Hash Tables
  - Linear Hash Tables
- ② Tree-based Index Structures
  - B+ Trees
- ③ Log-Structured Merge-Trees (LSM-Trees)


---

<sup>1</sup>Updated on April 12, 2021

## Hash-based Index Structures

### Secondary-Storage Hash Tables (外存哈希表)

- A secondary-storage hash table consists of a number of buckets
- An index entry with key  $K$  is put in the bucket numbered  $hash(K)$ , where  $hash$  is a hash function
- Each bucket stores a pointer to a linked list of pages holding the index entries in the bucket



# Categories of Secondary-Storage Hash Tables

## Static Hash Tables (静态哈希表)

- The number of buckets does not change

## Dynamic Hash Tables (动态哈希表)

- The number of buckets is allowed to vary so that there is about one block per bucket
- Extensible hash tables (可扩展哈希表)
- Linear hash tables (线性哈希表)

## Hash-based Index Structures Extensible Hash Tables


## Extensible Hash Tables (可扩展哈希表)

An extensible hash table is comprised of  $2^i$  buckets

- $i$  is called the global depth
- An index entry with key  $K$  belongs to the bucket numbered by the first  $i$  bits of  $hash(K)$

Example:


$hash(a) = 1010$ ,  $hash(b) = 0001$ ,  $hash(c) = 1001$ ,  $hash(d) = 1100$



## Extensible Hash Tables (Cont'd)


Every bucket keeps a pointer to a page where the index entries in the bucket are stored

- Several buckets can share a page if all the index entries in those buckets can fit in the page
- Every page records # bits of  $hash(K)$  (local depth) used to determine the membership of index entries in this page



## Properties of Extensible Hash Tables

- $\# \text{buckets} = 2^{\text{global\_depth}}$
- The global depth must be greater than or equal to the local depth of any page
- The page that a bucket points to is shared by another bucket if and only if the local depth of the page is less than the global depth




## Extensible Hash Table Lookup

Find the index entry with key  $K$

- ① Determine the bucket where the entry belongs to
- ② Find the entry in the page that the bucket points to

Example:  $K = a$ ,  $\text{hash}(a) = 1010$




## Extensible Hash Table Insert

Insert an index entry with key  $K$

- ① Find the page  $P$  where the entry is to be inserted
- ② If  $P$  has enough space, done!  
Otherwise, split  $P$  into  $P$  and a new page  $P'$

Example:  $K = f$ ,  $\text{hash}(f) = 0010$

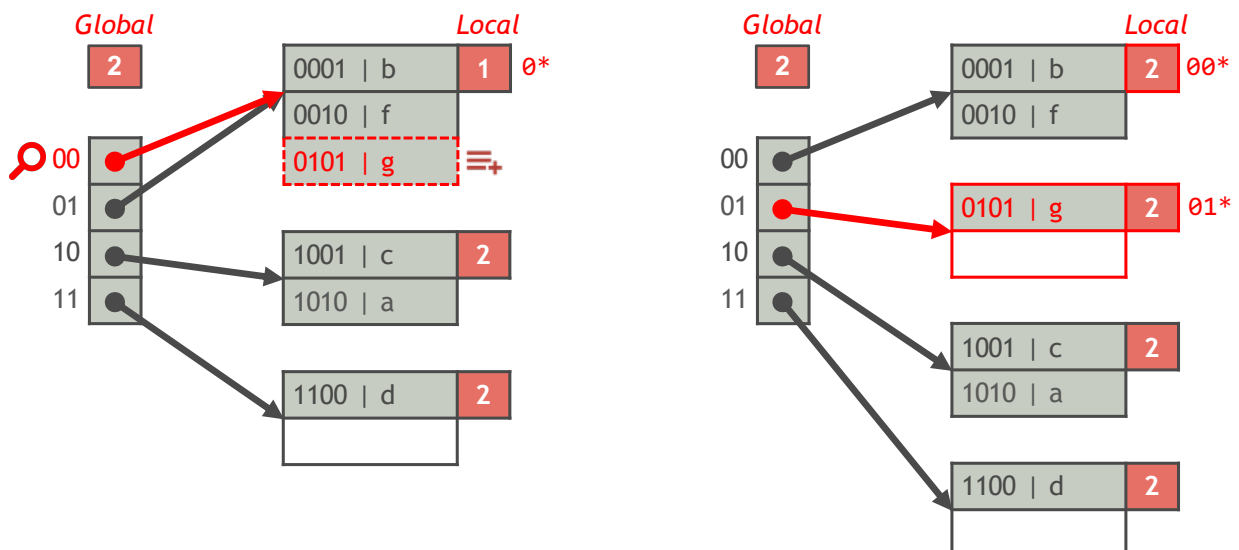


## Extensible Hash Table Insert (Cont'd)

If  $P$  overflows and the local depth of  $P$  is less than the global depth,

- ① Increase  $P$ 's local depth by 1
- ② Re-assign some index entries in  $P$  to a new bucket page  $P'$  ( $P$  and  $P'$  have the same local depth)

Example:  $K = g$ ,  $\text{hash}(g) = 0101$




## Extensible Hash Table Insert (Cont'd)

If  $P$  overflows and the local depth of  $P$  is equal to the global depth,

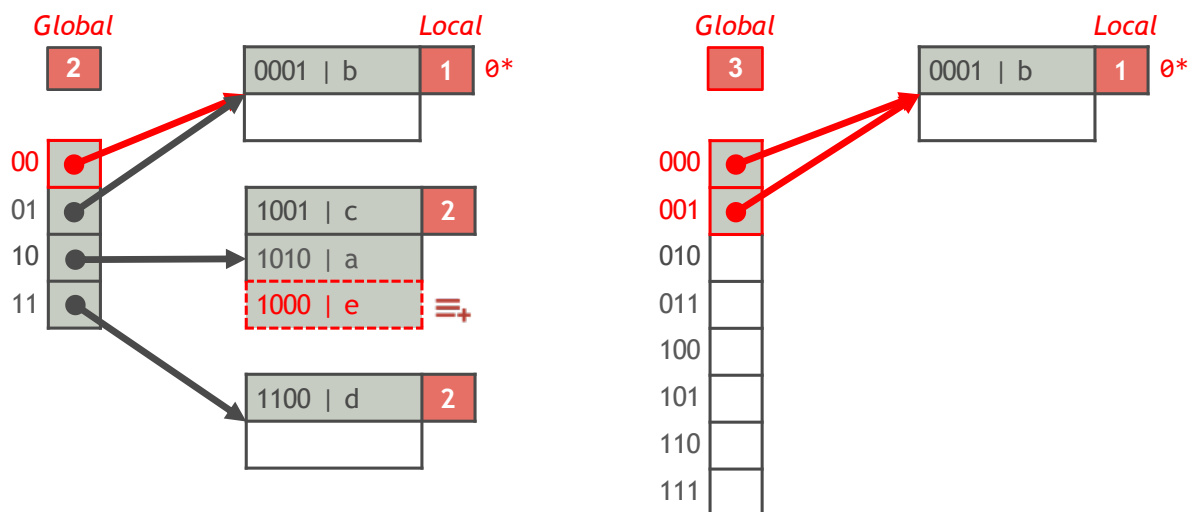
- ① Increase the global depth by 1 (double # buckets)
- ② Re-organize the buckets; if a page overflows, split it

Example:  $K = e$ ,  $\text{hash}(e) = 1000$





## Extensible Hash Table Insert: Example

Example:  $K = e$ ,  $\text{hash}(e) = 1000$



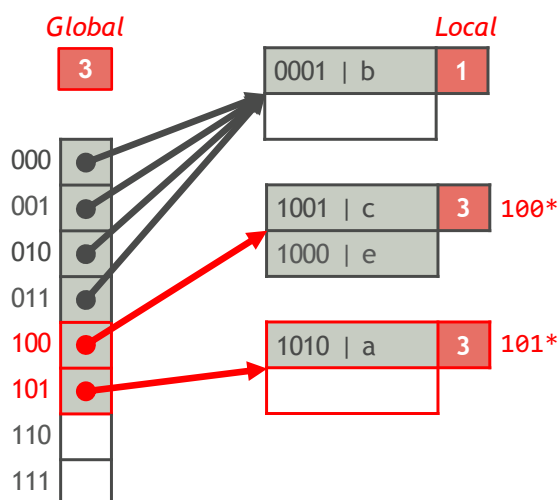

## Extensible Hash Table Insert: Example

Example:  $K = e$ ,  $\text{hash}(e) = 1000$



## Extensible Hash Table Insert: Example

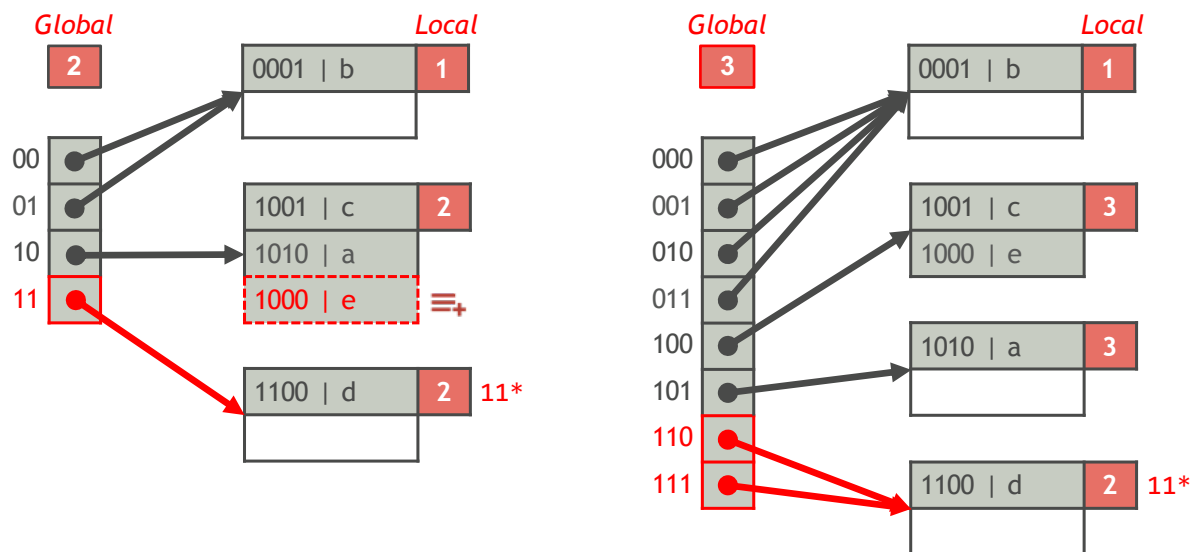
Example:  $K = e$ ,  $\text{hash}(e) = 1000$





## Extensible Hash Table Insert: Example

Example:  $K = e$ ,  $\text{hash}(e) = 1000$




## Extensible Hash Table Delete

Delete the index entry with key  $K$

- 1 Find the page where the entry belongs to
- 2 Delete the entry from the page

Example:  $K = a$ ,  $\text{hash}(a) = 1010$



# Hash-based Index Structures


## Linear Hash Tables

## Linear Hash Tables (线性哈希表)

A linear hash table is comprised of  $n$  buckets

- Every bucket keeps a pointer to a linked list of pages holding the index entries in the bucket
- Suppose each page can hold at most  $b$  index entries. The linear hash table stores at most  $\theta bn$  entries, where  $0 < \theta < 1$  is a threshold

Example:  $b = 2$ ,  $\theta = 0.85$



## Hashing Scheme

- The buckets are numbered from 0 to  $n - 1$
- Let  $m = 2^{\lfloor \log_2 n \rfloor}$ , so  $m \leq n < 2m$
- If  $\text{hash}(K) \bmod 2m < n$ , index entry with key  $K$  belongs to bucket  $\text{hash}(K) \bmod 2m$ ; Otherwise, it belongs to bucket  $\text{hash}(K) \bmod m$

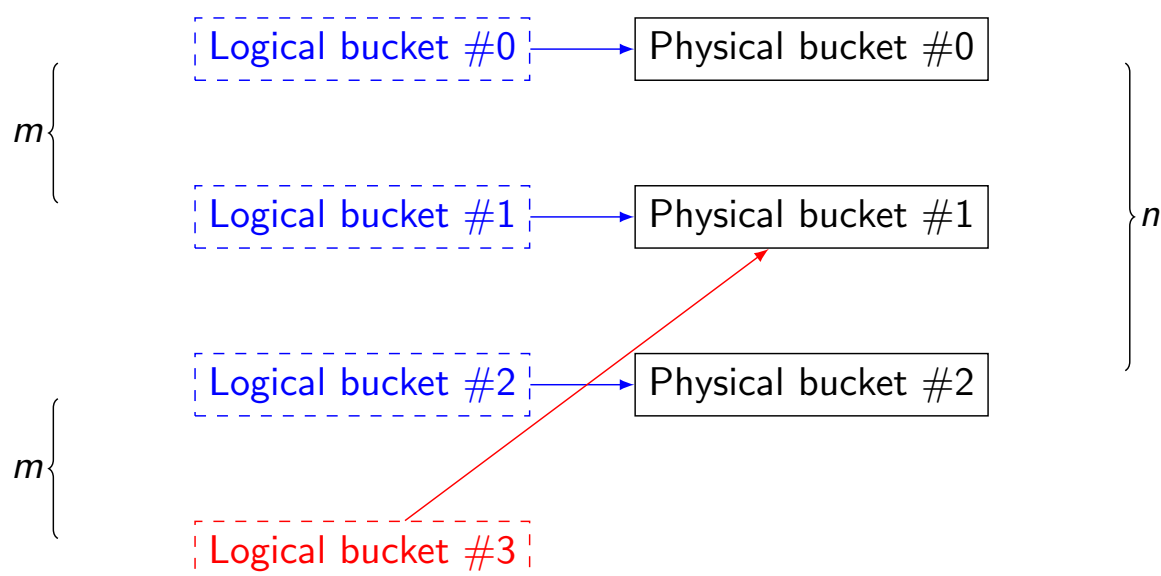
Example: Let  $n = 3$ . We have  $m = 2$  because  $2 = m \leq n < 2m = 4$

Bucket #0	$\text{hash}(K) = 0, 4, 8, \dots$
Bucket #1	$\text{hash}(K) = 1, \textcolor{red}{3}, 5, \textcolor{red}{7}, 9, \dots$
Bucket #2	$\text{hash}(K) = 2, 6, 10, \dots$

The buckets are NOT load-balanced

## Hashing Scheme (Cont'd)

- The logical bucket number  $b(K)$  for key  $K$  is  $\text{hash}(K) \bmod 2m$
- The physical bucket number for key  $K$  is  $b(K)$  if  $b(K) < n$
- The physical bucket number for key  $K$  is  $b(K) \bmod m$  if  $b(K) \geq n$




## Hashing Scheme (Cont'd)

- The buckets are numbered from 0 to  $n - 1$
- Let  $m = 2^{\lceil \log_2 n \rceil}$ , so  $m \leq n < 2m$
- If  $\text{hash}(K) \bmod 2m < n$ , index entry with key  $K$  belongs to bucket  $\text{hash}(K) \bmod 2m$ ; Otherwise, it belongs to bucket  $\text{hash}(K) \bmod m$

Example:

$\text{hash}(a) = 0101$ ,  $\text{hash}(b) = 0000$ ,  $\text{hash}(c) = 1111$ ,  $\text{hash}(d) = 1010$




## Linear Hash Table Insert

Insert an index entry with key  $K$

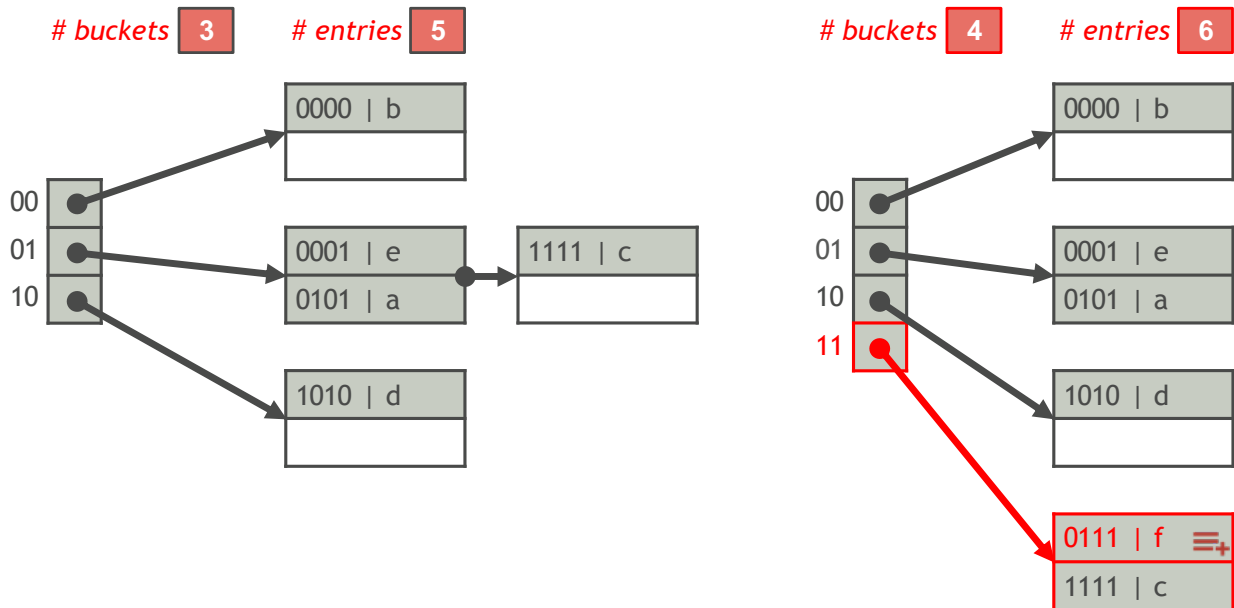
- 1 Insert the entry into the bucket  $B$  where it belongs to
- 2 Increase # entries by 1
- 3 If # entries  $\leq \theta bn$ , done!  
Otherwise, increase # buckets by 1 and redistribute the entries in  $B$

Example:  $\text{hash}(e) = 0001$ ,  $\theta = 0.85$



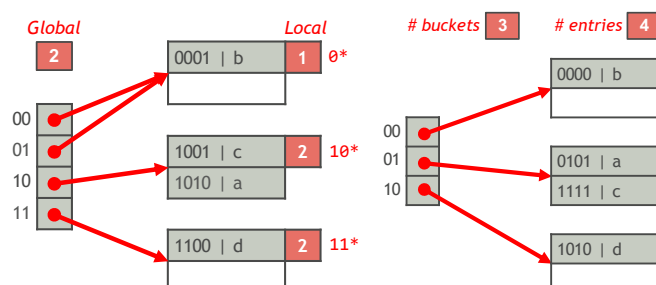
## Linear Hash Table Insert (Cont'd)

Example:  $hash(f) = 0111$ ,  $\theta = 0.85$



## Extensible Hash Tables VS Linear Hash Tables

	Extensible hash tables	Linear hash tables
# Buckets	$2^{global\_depth}$	$n$
Bucket pages	A bucket points to a single page	A bucket points to a linked list of pages
Hashing scheme	The first $global\_depth$ bits of $hash(K)$	$hash(K) \bmod 2^m$ or $hash(K) \bmod m$
Page split condition	A page overflows	$\#entries > \theta bn$
Hash table expansion	# buckets is doubled ( $global\_depth$ is increased by 1)	# buckets is increased by 1



# Tree-based Index Structures


# Tree-based Index Structures

## B+ Trees

## B+ Trees

A B+ tree is an  $M$ -way search tree with the following properties:

- It is perfectly balanced (i.e., every leaf node is at the same depth)
- Every node other than the root is at least half-full  
 $M/2 - 1 \leq \#keys \leq M - 1^2$
- Every inner node with  $k$  keys has  $k + 1$  non-null children
- Every node fits a page




<sup>2</sup>Raghu Ramakrishnan, Johannes Gehrke. Database Management Systems, 3rd Edition. 2003.

Navigation icons: back, forward, search, etc.

## B+ Tree Leaf Nodes

Every leaf node is comprised of an array of index entries (key/value pairs) and a pointer to its right sibling

- The index entry array is (usually) kept in sorted key order




Navigation icons: back, forward, search, etc.

## B+ Tree Inner Nodes


Every inner node is comprised of an array of keys and an array of pointers to its children

- The keys are derived from the attribute(s) that the index is based on
- The arrays are (usually) kept in sorted key order



## B+ Tree Root Node

The root contains at least one key






## B+ Tree Lookup

Find the index entry with key  $K$

- 1 Find the leaf node where  $K$  belongs to by following the direction of the keys in the inner nodes
- 2 Find the entry with key  $K$  in the leaf node

Example:  $K = 37$




## B+ Tree Range Query

Find the index entries with keys  $K \in [L, U]$


- 1 Find the first index entry  $E$  with the smallest key  $\geq L$
- 2 Scan the contiguous index entries with keys  $\leq U$  to the right of  $E$

Example:  $K \in [10, 25]$



## B+ Tree Range Query (Cont'd)

Example:  $K \in [10, 25]$



## B+ Tree Insert

Insert an index entry with key  $K$


- ① Find the correct leaf node  $L$  where the entry is to be inserted
- ② Put the entry into  $L$  in sorted key order
- ③ If  $L$  has enough space, done!  
Otherwise, split the keys in  $L$  into  $L$  and a new node  $L_2$ 
  - ① Redistribute the entries evenly, copy up the middle key
  - ② Insert an index entry pointing to  $L_2$  into the parent of  $L$

To split an inner node,

- ① Redistribute the entries evenly
- ② Push up the middle key


## B+ Tree Insert: Example 1 (w/o Node Split)

Example:  $K = 10$




## B+ Tree Insert: Example 1 (w/o Node Split)

Example:  $K = 10$




## B+ Tree Insert: Example 2 (w/ Node Split)

Example:  $K = 15$




## B+ Tree Insert: Example 2 (w/ Node Split)

Example:  $K = 15$




## B+ Tree Insert: Example 2 (w/ Node Split)

Example:  $K = 15$




## B+ Tree Insert: Example 2 (w/ Node Split)

Example:  $K = 15$




## B+ Tree Insert: Example 2 (w/ Node Split)

Example:  $K = 15$




## B+ Tree Insert: Example 2 (w/ Node Split)

Example:  $K = 15$



## B+ Tree Insert: Example 2 (w/ Node Split)

Example:  $K = 15$



## B+ Tree Delete


Delete an index entry with key  $K$

- ① Find the leaf node  $L$  where the entry belongs to
- ② Remove the entry from  $L$
- ③ If  $L$  is at least half-full, done!  
Otherwise,
  - ① Try to redistribute, borrowing from sibling
  - ② If redistribution fails, merge  $L$  and its sibling

If merge occurred, must delete entry pointing to  $L$  or the sibling from the parent of  $L$


## B+ Tree Delete: Example 1 (w/o Node Underflow)

Example:  $K = 17$



## B+ Tree Delete: Example 1 (w/o Node Underflow)


Example:  $K = 17$






## B+ Tree Delete: Example 2 (Key Redistribution)

Example:  $K = 7$




## B+ Tree Delete: Example 2 (Key Redistribution)

Example:  $K = 7$




## B+ Tree Delete: Example 2 (Key Redistribution)

Example:  $K = 7$




## B+ Tree Delete: Example 3 (w/ Node Merge)

Example:  $K = 11$




## B+ Tree Delete: Example 3 (w/ Node Merge)

Example:  $K = 11$




## B+ Tree Delete: Example 3 (w/ Node Merge)

Example:  $K = 11$



## B+ Tree Delete: Example 3 (w/ Node Merge)

Example:  $K = 11$



## B+ Tree Demo

<https://cmudb.io/btree>

## Key Compression

- The number of disk I/Os to retrieve a data entry in a B+ tree = the height of the tree  $\approx \log_{fan\_out}(\# \text{ of data entries})$
- The **fan-out** (扇出) of the tree is the number of index entries fit on a page, which is determined by the size of index entries
- The size of an index entry depends primarily on the size of the search key value
- Search key values are very long  $\implies$  the fan-out is low  $\implies$  the tree is high  $\implies$  the query time is long

## Prefix Compression (前缀压缩)

- Sorted keys in the same leaf node are likely to have the same prefix
- Instead of storing the entire key each time, extract common prefix and store only unique suffix for each key

Microphone	Microsoft	Microwave
------------	-----------	-----------

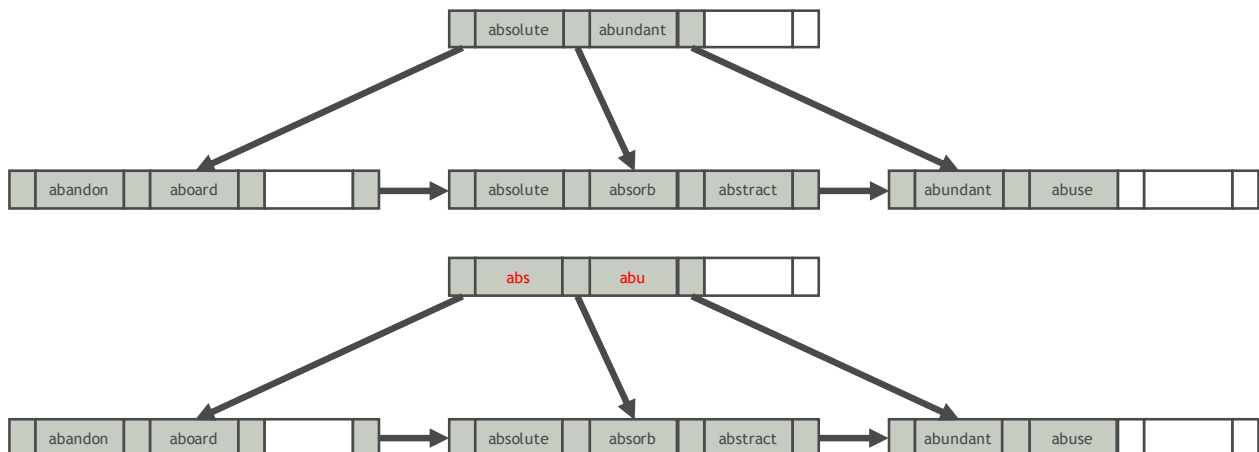
↓ Prefix compression

Prefix: Micro

phone	soft	wave
-------	------	------

## Suffix Truncation (后缀截断)

- The keys in the inner nodes are only used to direct traffic
- We need not store the keys in their entirety in inner nodes
- Store a minimum prefix that is needed to correctly route probes



## Bulk Loading (批量加载)

Creating a B+ tree on an existing set of index entries

### Top-Down Approach


- Insert the index entries one at a time
- Expensive, because each entry requires to start from the root and go down to the appropriate leaf node

### Bottom-Up Approach

- 1 Sort the index entries according to the search key
- 2 Allocate an empty inner node as the root and insert a pointer to the first page of sorted entries into it
- 3 Entries for the leaf pages are always inserted into **the right-most inner node** just above the leaf level. When that page fills up, it is split


## Bulk Loading: Example

Sorted keys: 2, 3, 7, 11, 17, 19, 23, 29, 31, 37




## Bulk Loading: Example

Sorted keys: 2, 3, 7, 11, 17, 19, 23, 29, 31, 37



## Bulk Loading: Example

Sorted keys: 2, 3, 7, 11, 17, 19, 23, 29, 31, 37



## Log-Structured Merge-Trees (LSM-Trees)



# Log-Structured Merge-Trees (LSM-Trees)

The log-structured merge-tree (LSM-tree) has been widely adopted in the storage layers of modern NoSQL systems

- LevelDB, RocksDB, HBase, Cassandra, TiDB and so on

An LSM-tree applies **out-of-place updates**

- First, it buffers all writes in memory
- Subsequently, it flushes the writes to disk and merges them using sequential I/Os

## Basic LSM-Tree Structure

- **Memtable**: a mutable B+ tree or hash table in main memory
- **Immutable file**: an immutable sorted file on disk

**Before compaction (merge)**

Memtable (3, 333), (7, 777) *in memory*

Immutable file (2, 222), (3, 123), (5, 555), (8, 888) *on disk*

**After compaction**


Memtable Empty

Immutable file (2, 222), (3, 333), (5, 555), (7, 777), (8, 888)

**Disadvantage:** When the immutable file is very large, the lookup and the merge on the file are costly

## Leveled LSM-Tree Structure

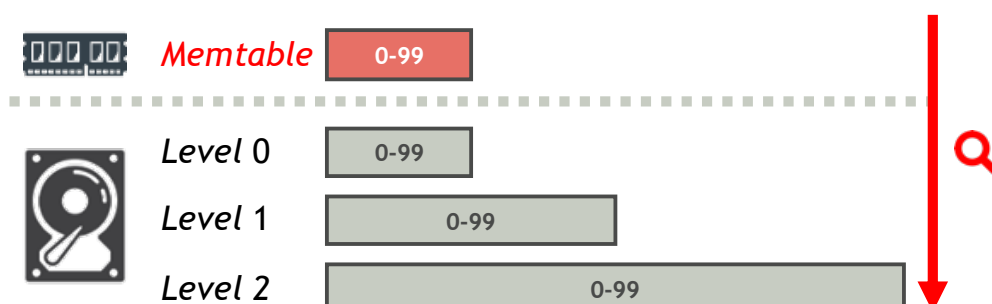
- **Memtable**: a mutable B+ tree or hash table in main memory
- **Level 0**: an immutable copy of the memtable on disk
- **Level  $i$  ( $i \geq 1$ )**: immutable sorted file on disk
  - ▶ The key-value pairs in level  $i + 1$  are older than those in level  $i$
  - ▶ Level  $i + 1$  is  $T$  times larger than level  $i$



## LSM-Tree Lookup

**Find the latest value for the key  $K$**


- 1 Find the key-value pair with key  $K$  in the memtable. If  $K$  is contained in the memtable, return the value for  $K$  (if a tombstone of  $K$  is found, return "not found"); Otherwise, find the latest value for the key  $K$  in the levels.
- 2 While  $i \leq n$ , find the key  $K$  in level  $i$ . If  $K$  is contained in level  $i$ , return the value for  $K$  (if a tombstone of  $K$  is found, return "not found")
- 3 If  $K$  is not contained in level  $n$ , return "not found"



# LSM-Tree Insertion

## Insert a key-value pair ( $K, V$ )

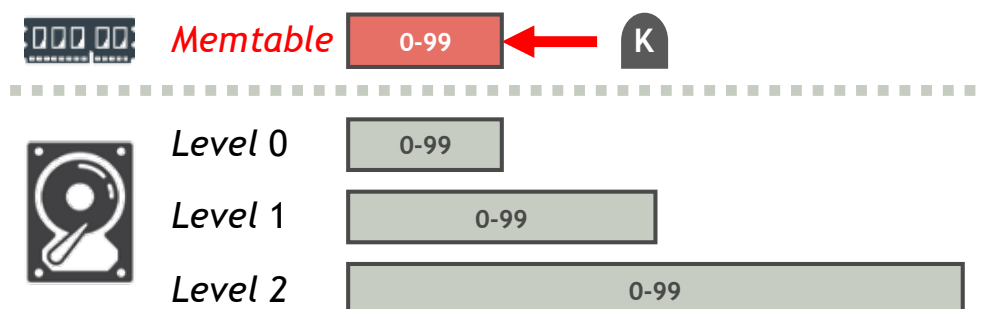
- 1 Insert ( $K, V$ ) into the memtable (in-place update)
- 2 If the memtable does not overflow (exceed its maximum allowable size), done! Otherwise, compact the key-value pairs in the memtable to level 0
- 3 If level  $i$  overflows, compact the key-value pairs in level  $i$  to level  $i + 1$



# LSM-Tree Deletion

## Delete the key-value pair with key $K$

- 1 Insert a *tombstone* for key  $K$  into the memtable
- 2 In compaction, any elder key-value pair with key  $K$  are deleted when it is merged with the tombstone of  $K$



## B+ Trees VS LSM-Trees

	B+ trees	LSM-trees
Update method	In-place update	Out-of-place update
Space amplification	Low (only one copy for a key)	High (many copies for a key)
Write performance	Low (random I/Os)	High (sequential I/Os)
Space utilization	Fragmentation (1/4 of a page is not used)	High (key-value pairs are compacted into sorted runs)
Concurrency control and failure recovery	Complicated	Simple (sorted runs are immutable, and compaction is out-of-place)

## Summary

- 1 Hash-based Index Structures
  - Extensible Hash Tables
  - Linear Hash Tables
- 2 Tree-based Index Structures
  - B+ Trees
- 3 Log-Structured Merge-Trees (LSM-Trees)

- ① 当B+树进行删除操作时，若一个节点不足半满，是优先向左兄弟借，还是优先向右兄弟借呢？

答：都可以，取决于B+树的具体实现方法。